

POLITECHNIKA BIAŁOSTOCKA

WYDZIAŁ INFORMATYKI

PRACA DYPLOMOWA INŻYNIERSKA

TEMAT: SKELETAL ANIMATION USING
INVERSE KINEMATICS IN THE UNITY
ENGINE

WYKONAWCA: ŁUKASZ BIAŁCZAK

.....
podpis

PROMOTOR: DR INŻ. ADAM BOROWICZ

BIAŁYSTOK 2022 r.

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Problem Formulation	2
2	Related Work	3
2.1	IK Algorithms	3
3	Tools	4
3.1	MakeHuman	4
3.2	Unity	4
3.2.1	Importing Animations	4
3.2.2	Animator Controller	4
3.2.3	Animation Rigging Package	4
4	Inverse Kinematics in the Unity Engine	5
4.1	FABRIK implementation	5
4.2	Spider Movement	7
4.2.1	Project Setup	7
4.2.2	Scripts	8
4.3	Human Animation Sequence	11
4.3.1	Project Setup	11
4.3.2	Scripts	11
5	Experiments	12
6	Conclusion	13

1. Introduction

1.1 Motivation

Animation is the technique of displaying different positions of a character or object in rapid succession to create the illusion of movement. It is used in various forms of entertainment, such as movies and video games. In the latter, unlike in the former, the animation sequences are performed in real time and therefore impose additional constraints. Without the freedom to process a single frame for minutes or hours during the rendering of the scene, the animator must compromise on the quality and realism of the sequence in order to optimize for gameplay. One such optimization is the use of skeletal animation in which animation sequences are performed by manipulating a tree-like structure of interconnected bones, represented by transforms, to create the desired motion of the character. Furthermore, the interactive nature of video games makes it impossible for the artist to create predefined animation sequences for every possible situation that may occur in the game. For example, an animation of a hand pressing a button is only valid if the character to which the hand belongs is placed exactly in a predefined position. If the button changes its size or position, the baked animation must be altered, or it simply loses its realism. As a result, predefined animation sequences are often generic and do not allow the character or object to interact naturally with their surroundings. Game developers have come up with many methods to improve the realism of animation in games such as playing cutscenes for critical interactions between a character and the world. However, this paper will focus on the use of procedural animation and, more specifically, the application of inverse kinematics to skeletal animations in video games.

Inverse kinematics is a technique used in fields such as robotics and computer graphics to determine the joint angles of a kinematic chain that will result in a particular part of the chain, usually an end effector, reaching a specified position in 3D space. In computer graphics specifically, the technique is often used to animate the movement of characters and objects such that they interact with their surroundings in a more realistic manner. Inverse kinematics is even used in modeling and animation software such as Blender [2] to speed up the process of limb manipulation when creating a static animation.

There are multiple approaches and algorithms that exist within the inverse kinematics domain, such as analytical methods, gradient descent, and optimization techniques. The choice of approach varies depending on the complexity of the use case, the desired realism of the animation, and system limitations.

1.2 Problem Formulation

The aim of this dissertation is to gain a better understanding of the basic algorithms used in inverse kinematics, discover the built-in functionalities that the Unity engine offers for such implementation. The project implementation will apply these concepts to create pairs of animations which consist of baked and inverse kinematics variants. The use cases will expand the problem by introducing additional constraints which will be required to keep the consistency and realism of the animations. The variations will then be compared through the lens of realism and performance.

The author will begin by discussing the theory of the different approaches and algorithms used to solve the inverse kinematics problem, and the resulting choice of the algorithm to be used in the project implementation. The following sections will explain in depth the implementation of two use cases which demonstrate the purpose of inverse kinematics as a skeletal animation technique. Experiments will then be conducted to compare the inverse kinematics animations with their baked counterparts based on realism and performance. Finally, a summary and conclusion of important points will be presented to the reader.

2. Related Work

Inverse kinematics is applied in various fields such as robotics and computer graphics. Over time, many approaches have surfaced in order to solve the IK problem. There are multiple families of solutions [1] which suit different use cases.

2.1 IK Algorithms

3. Tools

Multiple tools were used in the process of creating the demo application for this paper including the Unity game engine, Blender as a modeling and animation software, and MakeHuman as a model creation tool. This chapter discusses the built-in functionalities which make the mentioned tools an effective choice.

3.1 MakeHuman

MakeHuman is an open source tool for making 3D characters. It provides a convenient way of acquiring a human model which is customizable and can be exported in various formats in order to be used in other softwares. The key factors which make this tool suitable for use in the demo application is the options it provides regarding the complexity of the topology of the model's mesh, and the choice of skeleton rig alongside the fact that the exported model is already rigged and ready to be used in an animation software. One of the rig options is specifically designed to be then used in a game engine setting (See Figure).

3.2 Unity

The Unity game engine is the one tool which was non-negotiable as the paper is meant to specifically focus on the usage of inverse kinematics in said game engine. Nevertheless, the engine is a good selection for this use case due to its advanced 3D support, the built-in packages and functionalities which are geared towards the subject of this paper, and the overall popularity of the engine and large community built around it which results in a substantial amount of documentation and support.

3.2.1 Importing Animations

3.2.2 Animator Controller

3.2.3 Animation Rigging Package

4. Inverse Kinematics in the Unity Engine

The demo application written for the purpose of this dissertation includes two separate use cases of skeletal animation using inverse kinematics in the Unity engine. The first example is that of a four legged spider which uses the technique as a means to more naturally adjust its limbs to the terrain it moves around upon. The second example is the application of inverse kinematics to an animation sequence of a human character pressing multiple buttons in succession. The use of inverse kinematics allows the character to adjust its animation to hit all the buttons without the need for a baked animation targeted towards each button, as well as dynamically adjust the order of the buttons to be hit.

Although there are two separate use cases demonstrated in this application, both use the same implementation of the FABRIK algorithm.

4.1 FABRIK implementation

This implementation of the FABRIK algorithm is based on the paper written by the author of the Unity engine FABRIK implementation. (CITE THE PAPER LMAO). For the purposes of this application, the basic algorithm is implemented without many additional constraints and options available. The one constraint added on to this implementation is that of pole targets which will be further explained when discussing the code behind them.

The script which implements the algorithm in this project takes in a few parameters required to set up the mechanism which are shown in figure 4.1. First and foremost, the root and leaf nodes must be provided in order to define the kinematic chain which is to be manipulated. The next object which the script must have knowledge of is the target transform which the end effector will attempt to move to. The script must also have a tolerance parameter which dictates how close the end effector must be to the target for the position to be considered as a valid solved position. All the aforementioned parameters are required for the script to function. Additionally, a pole target object may be optionally passed in if the use case requires it to function in a desired manner.

The first case which the algorithm must cover is if the distance from the root to the target object is greater than the sum of distances between each adjacent bone in the defined

Chain Root	Bone_L.001 (Transform)	⊙
Chain Leaf	Bone_L.003_end (Transform)	⊙
Target	fl_target (Transform)	⊙
Pole	fl_pole (Transform)	⊙
Tolerance	0.001	

Figure 4.1: FABRIK script parameters

kinematic chain. In this case, the target is out of reach. Given that the bones in a skeleton are expected to keep a fixed length, the end effector will not be able to reach the target, and instead the kinematic chain straightens and extends in the direction of the target. There is no need to continue with the iterative portion of the algorithm.

In the implementation below, square magnitudes are used to avoid the calculation of square roots, thus slightly optimizing the initial check.

The scenario where the target is within the reach of the kinematic chain utilizes iterative forward and backward component after which the algorithm is named. Before each iteration, the positions of the joint transforms are copied. All operations and calculations are performed on these copied transforms and at the end of the full pass, the new positions are then copied back to the kinematic chain. **CHECK IF THE ALGORITHM WORKS WITHOUT THIS. MIGHT SAVE SOME MEMORY LMAO.** It is also important to note that the modification of the transforms is done in Unity's *LateUpdate* function. When using a mix of inverse kinematics and baked animation, the object to which the IK script is attached will have Unity's built-in *Animator* component. If the IK script updates joint transforms in the *Update* method, then they may be overwritten by the *Animator* component.

The forward pass of the algorithm iterates through the chain starting from the end effector and ending at the root. At the start, the end effector's position is set to be equal to the position of the target. A straight line is then drawn from the end effector to the following node. This neighbour's position is then interpolated along the line so that the original distance between the two nodes is kept. The same operation is performed for each pair of neighbouring nodes throughout the pass.

When the forward pass is complete, the root node is displaced from its original position. This is undesired, as the root's node position should not be affected by the algorithm.

To remedy this, the next step is to repeat the forwards pass, but this time in reverse. The root node's position is set equal to what it was at the beginning of the frame. The next node is then interpolated between its current position and the root to keep the initial bone length. As with the forward pass, this is repeated for each subsequent pair of nodes.

These two steps are repeated together until the end effector is within a threshold distance of the target. The FABRIK algorithm is a heuristic algorithm, and as such it does not lead to an exact result. Instead, it aims to approximate the correct solution and solves the problem in a less complex and more optimized way. Again, the square distances are used to avoid the calculation of square roots.

POLE TARGETS. FIND A SOURCE WHICH EXPLAINS THE METHOD USED IN MY PROGRAM. IF NOT THEN REFER TO THE YOUTUBE VIDEO LOLOL

4.2 Spider Movement

The first use case for inverse kinematics in the demo application is that of a four legged spider. The algorithm is used to adjust the creature's limbs to uneven terrain, leading to a much more natural and realistic movement.

4.2.1 Project Setup

Each one of the spiders legs is treated as a separate kinematic chain. The spider prefab consists of a container which holds the spider object itself, set of empty objects to which the four IK scripts are attached for the purpose of easy yet separated access. The prefab also contains sets of ray casts and targets. The ray casts serve to scan the surface of the terrain under the spider, and mark the targets to which each leg should move.

Ray casts are dispatched from above the spiders legs, and aim in the creatures local negative Y axis. This ensures that no matter what orientation the spider finds itself in, the rays are always pointing at the surface which it is standing on. Masks are applied to the rays, making sure that only terrain objects are taken into account, while the creature's body itself is not. The ray cast hit point positions are then applied to each leg's respective target object. These targets serve as markers for the limbs end effectors.

4.2.2 Scripts

With the project set up in this manner, scripts must now be added on to make the scene functional. A script is required for the main movement of the spider, the ray cast logic, and the mechanism which controls the ik targets.

Ray casts

The ray cast objects contain a script component which dispatches the rays and sets the appropriate target positions. First, a mask must be established, which will then be passed into the ray cast operation. This is required so that only terrain is counted as a valid hit. The lack of such a mask may result in unexpected behavior, such as targets being set on the spider's body itself. The ray cast object is then created, shooting in the local negative Y axis direction. This ensures that no matter the orientation of the creature, the rays are sent towards the surface that the spider is standing on. The targets controlling the spider's legs will not be updating their positions to the ray cast hit points each frame, though their scripts must have knowledge of the hit positions at any given time. Given this, a separate set of objects are set to track the ray cast hit points each frame.

Target Logic

In order for the spider's movement to seem realistic, the targets controlling each leg must adhere to a set of rules pertaining to their movement. As mentioned in the ray cast section, the IK targets cannot simply be set to track the ray cast hit points. The following is an outline of the rules specified for the IK targets, which define if it should start moving towards its ray cast hit target:

- A target must be grounded to be eligible for a movement sequence.
- A target will only begin moving towards the ray cast hit point if the distance between them is above a specified threshold.
- A target is only allowed to begin moving towards the ray cast hit point if both legs on the opposite diagonal are grounded (See Figure 4.2).




Figure 4.2: The diagonals referencing the spiders legs which are used when checking if a target is allowed to begin moving, where A and D are on an opposite diagonal to B and C

When a target satisfies all of these conditions, it makes note of the ray cast hit's current position, which it will use for its upcoming movement sequence. Once it begins moving, the *Grounded* boolean is set to be false, so that the other targets know whether they can begin moving or not.

The movement sequence itself is done using Unity's *Vector3.MoveTowards* method, which takes in a current position vector, a vector to move to, and the maximum distance to move per frame, which can be used to control the movement speed. This method allows the target to interpolate its position every frame. The values fed into this method are simply the target's current position, the ray cast hit target's position, which was recorded right before the beginning of the movement sequence, and an arbitrary speed value, which is dependent on *Time.deltaTime* to avoid variations when the frame rate changes. The only caveat is that in the first half of the movement, the destination vector's height component is increased to achieve an arc-like movement. This produces the effect of lifting the spider's leg.

General Movement

The main movement script for the IK spider implementation brings the whole system together. It has three main objectives:

- Calculate the rotation of the spider so that it's local up and forward vectors can be set

accordingly.

- React to input by moving the main body of the spider along with the ray casts.
- Regulate the height of the spiders body above the ground.

The first objective - calculating the rotation of the spider - is what allows it to scale walls and walk upside down on the ceiling. The rotation is determined based on the limb positions at the beginning of each *Update* call. First, two vectors are constructed from the end effectors of both sets of diagonally opposed legs. The local up vector of the spider is then calculated by taking the cross product of these two vectors. This determines the orientation of the spider's main body, and it also affects the direction that the four legs' rays are cast. The local forward vector is then obtained from the cross product between the newfound up vector and the spider's right vector. This new vector is what will be used to determine the direction of movement when the spider receives input from the user.

The second objective - reacting according to input - is quite simple once the local directional axes are determined. When the script detects a non-zero value on either the *Horizontal* or *Vertical* axis, it reacts accordingly by moving the spider along with all the ray cast objects. The *Vertical* axis corresponds to the forward and backward movement of the spider. The script reacts by simply moving the spider's main body along the local forward axis mentioned previously. Additionally, the ray casting objects are offset either forwards or backward depending on the direction of the movement. This is done because the ray casts must be slightly ahead of the default leg positions so that the legs end up moving in a natural manner. The *Horizontal* axis is responsible for rotating the spider about its local up axis which allows the spider to turn around.

Finally, the spider's height off of the surface must be regulated each frame. The lack of a gravitational force acting on the body to keep it flush with the ground, and the unconstrained rotational capability of the creature, means that the distance between the spider and the surface it is walking on must be procedurally kept in check. This is done with yet another ray cast which originates from the center of the spider's body and points in the negative up axis direction. If the distance to the hit point exceeds an acceptable range, the body is moved towards said range, again using the *Vector3.MoveTowards* method to linearly interpo-

late the spiders position and avoid excessive jerkiness which occurs with a frequent variation in height.

4.3 Human Animation Sequence

4.3.1 Project Setup

4.3.2 Scripts

5. Experiments

6. Conclusion

Bibliography

- [1] Andreas Aristidou and Joan Lasenby. Fabrik: A fast, iterative solver for the inverse kinematics problem. *Graphical Models*, 73:243–260, 9 2011.
- [2] Blender 3.4 Manual. Inverse kinematics introduction. https://docs.blender.org/manual/en/latest/animation/armatures/posing/bone_constraints/inverse_kinematics/introduction.html, 2022. Accessed: 2022-12-26.